


a

Disagreement-based preference mining

Open-ended tasks



Existing datasets New outputs

Metrics and judges

Signal 1 Signal 2 Priority
 $A > B = A > B$ Lower

$A > B \neq B > A$ Higher

Conflict → rank higher
 + random sample



Reuse labels
or ask humans
for new ones

Human feedback

Signals can share errors

b

Depth-first evaluator evolution

Rejection alone does not end a direction



Active direction: make layout errors count

Target error: layout faults overlooked

Base model fixed

What should we test next?

Candidate



Add a layout
detector

Execute and inspect

Checking result ↓

Keep evidence after rollback



Retain diagnostics

Observed

Layout errors detected.
Final scores unchanged.

Question to test

Do findings reach scoring?

Next revision



Pass findings
into scoring

Same direction

End with reason → new direction



Evaluate every candidate

Which program
should we keep?



Same checking set

Improves + passes checks
→ update best program

No gain → keep best program




Best program

c

Downstream applications

Output selection



Rank and select

Critic-guided
refinement



Reward-guided
training



Reward
Planned