

a Task: evaluate edits

Make a rainy night.
Keep the street layout
and buildings unchanged.

Each edit is scored separately.

Source



A



B



Layout fidelity

Preserved

Wider street

Visual richness

Less vivid

More vivid

Human: A > B


≠

Evaluator: B > A

b Default setup

Try a different direction

Initial evaluation program



Compare layouts



Change found;
not scored.


× B > A

Check night scene



Rain at night?
Wet pavement?



Inspect details



Windows?
Reflections?




Direction left undeveloped



c IterEval: depth-first evolution

Continue one direction: layout fidelity

Initial evaluation program




1 Detect layout changes



Detected, but
not scored.

× B > A

2 Use changes in scoring



Scored, but
underweighted.

× B > A

3 Weight by severity



Larger changes,
larger penalties.

✓ A > B

Local ordering ≠ candidate acceptance.

Illustrative search paths. Nodes are complete evaluation programs; images stay fixed.