

a Task: evaluate image edits

Make a rainy night; keep the street layout and buildings unchanged.

Source

A

B



Layout fidelity

Preserved

Wider street

Visual richness

Less vivid

More vivid



Human: A > B

≠




Evaluator: B > A

Each edit is scored separately.

b Default setup


Try a different direction



c IterEval: depth-first evaluator evolution

Continue one direction: layout fidelity

Candidate acceptance and program inheritance are checked separately.
Initial evaluation program



1 Detect layout changes



Changes detected, but unused in scoring.

× Still favors B

2 Penalize layout changes



Penalty too small; error remains.

× Still favors B

3 Scale the penalty by severity



Larger layout changes, larger score penalties.

✓ Now favors A

Illustrative search paths. Nodes are complete evaluation programs; images stay fixed.