

EVOLVING EVALUATION SYSTEMS FROM HUMAN FEEDBACK FOR OPEN-ENDED TASKS

Anonymous authors

Paper under double-blind review

Working manuscript: quantitative displays use simulated data; empirical evaluation pending.

ABSTRACT

Open-ended evaluation is difficult because quality depends on task-specific criteria and trade-offs. Human preferences provide judgments on concrete outputs, but do not specify what an evaluator should inspect or how its findings should affect future judgments. We study how an evaluation agent can evolve an evaluation procedure from human feedback. Doing so requires acquiring informative preferences and turning the problems they reveal into working changes to the evaluation procedure. A limited annotation budget can be spent on repetitive comparisons that provide little guidance for improving the evaluator. Disagreement-based preference acquisition allocates fewer annotations to comparisons on which existing evaluation signals agree. In the iterative optimizer we tested, rejected candidates were repeatedly followed by new proposals rather than further development of the rejected implementation. Yet a direction may require coordinated changes across the evaluation procedure before it produces a gain. We introduce ERA, falsification-guided depth-first evolution that separates candidate acceptance from direction continuation. A failed candidate does not by itself refute a direction’s hypothesis. ERA checks whether the proposed behavior occurred before treating its outcome as a test of the hypothesis. Diagnostic results guide further implementation or support refutation. Our cross-domain evaluation design tests acquisition and evolution separately and end-to-end under matched annotation and search budgets. Quantitative displays are simulated; confirmatory evaluation remains pending.

1 INTRODUCTION

Evaluating open-ended outputs is difficult because their quality depends on task-specific criteria and trade-offs. A shorter summary may preserve the central argument better than a longer one, while an elaborate webpage may still fail its main purpose. Metrics and model-based judges capture useful aspects of quality, yet do not consistently match human preferences (Fabbri et al., 2021; Xu et al., 2023; Zheng et al., 2023). Collecting more evaluation signals does not by itself resolve the central question: which evidence should matter, and how should conflicting criteria be weighed for this task?

Human preferences provide feedback on which output is better, but leave open what an evaluator should inspect and how its findings should affect future judgments. We study how an evaluation agent can evolve an evaluation procedure from human feedback by changing how it gathers evidence and makes judgments. Prior work uses human feedback to refine evaluation criteria or select evaluation functions (Liu et al., 2024; Shankar et al., 2024), while agentic judges gather evidence with tools (Zhuge et al., 2025). Methods that propose and test changes to prompts and agent code provide a starting point for evolving this procedure (Agrawal et al., 2026; Zhang et al., 2026). Evolving the procedure requires two linked steps: acquiring informative preferences and turning the problems they reveal into working changes.

Obtaining informative preferences is difficult in a new domain. Hand-designed comparisons can embed the designer’s assumptions; unrestricted collection can produce repetitive judgments. Our disagreement-based preference acquisition allocates fewer annotations to comparisons on which




Figure 1: Evaluation-agent self-evolution from human feedback. (a) Four artifact families retain distinct evaluation standards. (b) Reused and newly collected preferences enter a common evidence interface; their provenance and annotation quality remain separate. (c) ERA evolves a six-component program with a fixed base model. (d) C1 selection, C2 refinement, and planned C3 reward training. Artifacts and feedback are constructed illustrations; the photograph is generated with GPT Image 2. Arrows denote information flow, not measured improvements.

existing evaluation signals agree, drawing on query-by-committee and active preference learning (Seung et al., 1992; Muldrew et al., 2024). Random sampling checks for shared blind spots among these signals. We test whether these preferences support greater evaluator improvement than random sampling at equal human annotation cost.

Even with informative preferences, search can leave a direction unfinished. In the iterative optimizer we tested, candidates needed to improve the development score to replace the best evaluator. After rejections, search repeatedly returned to that evaluator and proposed other directions instead of further developing the rejected implementation. This pattern matters when a direction requires coordinated changes across the evaluation procedure. For example, evaluating omissions in a summary requires detecting missing content, judging its importance, and making that judgment affect the final score. Changing only one step may leave the overall judgment unchanged.

Rejecting a candidate need not end work on its direction. ERA separates *candidate acceptance* from *direction continuation* through falsification-guided depth-first evolution (Figure 1c). Each direction is motivated by a hypothesis about how a change would correct an evaluation error. A failed candidate does not by itself refute that hypothesis. ERA first checks whether the proposed behavior was implemented and affected the final judgment as hypothesized. New diagnostic results guide further work on the implementation. Results contradicting the hypothesis can end a direction; budget exhaustion alone does not refute it. We test whether this continuation produces more sustained gains than restarting proposals under the same search budget.

Preference acquisition and ERA target two obstacles to learning an evaluation procedure from human feedback: redundant annotation and premature abandonment of corrective directions. The cross-domain evaluation separates acquisition efficiency, search progress, and final evaluator agreement across structured visual artifacts, images and 3D, text generation, and automated research. It includes reused and newly collected preferences, with a pointwise evaluator for each task standard. The comparisons include evaluator fine-tuning (Kim et al., 2024a;b) and score calibration under controlled feedback and observation access.

2 RELATED WORK

Evaluators and feedback acquisition. G-Eval and MT-Bench study model-based judgment and human agreement (Liu et al., 2023; Zheng et al., 2023). Prometheus, Prometheus 2, and Self-Taught Evaluators improve judges through rubric-conditioned or synthetic training data (Kim et al., 2024a;b; Wang et al., 2024). AutoCalibrate refines evaluation criteria from human-labeled examples, while EvalGen helps users define criteria and select evaluation functions using human grades (Liu et al., 2024; Shankar et al., 2024). For DeepResearch reports, Lv et al. (2026) train query-specific rubric generators from human preferences and use the learned rubrics as training rewards. Our optimization target includes both the criteria and the procedure for gathering and using evaluation evidence.

RewardBench and JudgeBench assess evaluators on preference comparisons (Lambert et al., 2025; Tan et al., 2025). Benchmark² instead examines benchmark quality through ranking consistency, discriminability, and capability alignment, and studies selective test-set construction (Qian et al., 2026). Our acquisition question concerns which comparisons to annotate for evaluator improvement, drawing on query-by-committee and active preference learning (Seung et al., 1992; Muldrew et al., 2024). Because our evaluation signals can measure different properties, acquisition distinguishes conflicts within a dimension from trade-offs across dimensions and retains random sampling to check shared blind spots.

Optimizing language-model programs. DSPy exposes compositional pipelines to optimization (Khattab et al., 2024). ProTeGi, OPRO, TextGrad, and GEPA use model-generated feedback or proposals to improve instructions and programs (Pryzant et al., 2023; Yang et al., 2024; Yuksekgonul et al., 2025; Agrawal et al., 2026). ADAS, AFlow, and the Darwin Gödel Machine investigate automated construction or evolution of agentic systems (Hu et al., 2025; Zhang et al., 2025; 2026). GEPA retains candidates with complementary strengths, while DGM explores an archive of agent implementations. Studying the path from experience generation through skill extraction to consumption, Huang et al. (2026) report average gains alongside negative transfer and differences between a model’s ability to extract and use skills. Their study tests whether extracted skills help a target agent; ERA examines how to continue an evaluator-improvement direction after an unsuccessful implementation, separately from accepting that candidate.

Domain quality and downstream use. SummEval and BERTScore address summarization evaluation and text similarity (Fabbri et al., 2021; Zhang et al., 2020); ImageReward, Pick-a-Pic, and GenEval capture different aspects of image preference and alignment (Xu et al., 2023; Kirstain et al., 2023; Ghosh et al., 2023). These signals are useful evidence, not interchangeable definitions of quality. ReAct, Reflexion, and Self-Refine connect observations, reflection, and iterative improvement (Yao et al., 2023; Shinn et al., 2023; Madaan et al., 2023). Human-feedback summarization, instruction tuning, and direct preference optimization connect evaluation to model learning (Stiennon et al., 2020; Ouyang et al., 2022; Rafailov et al., 2023). The reward-hacking survey of Wang et al. (2026) organizes risks of proxy optimization around objective compression, optimization amplification, and evaluator-policy co-adaptation. An evaluator’s agreement on fixed outputs therefore needs to be distinguished from its behavior under optimization, motivating separate tests of fixed-pool selection, critic-guided refinement, and reward-guided training.

3 SCOPE AND PROBLEM FORMULATION

A *domain* is a collection of tasks sharing core quality dimensions, their priorities, the overall decision rule, and the human annotation protocol. An artifact type does not define a domain: screenshot-conditioned webpage reconstruction and specification-conditioned webpage generation require different standards. A change of source distribution that preserves the standard is instead a within-domain distribution shift. Appendix A enumerates the full research scope of 22 domains in four families, without assuming that all are ready for experiments. Closed tasks whose complete quality criterion is an executable postcondition fall outside this scope.

Existing human preferences are not a domain-admission requirement. Their availability determines the missing data work, not the algorithm’s definition. An evaluation ecosystem can exist without a

suitable preference dataset. We distinguish *reused preferences*, which test improvement on established benchmarks, from *new preferences*, which test applicability when matching human judgments must first be collected. One domain can support both settings.

Let x be a task with permitted context, y a candidate output, and $\mathcal{Q}_D = \{q_1, \dots, q_m\}$ the quality dimensions for domain D . A human comparison has the core form

$$h = (x, y_a, y_b, q, p, a, \nu), \quad (1)$$

where q is a dimension or overall quality, p is a preference, tie, or abstention, a records annotator agreement, and ν records provenance. Feedback has four levels: H_0 gives overall preference; H_1 gives dimension-wise preference; H_2 identifies error type, severity, and location; H_3 optionally supplies a free-form rationale. H_1 is central to the dimension-guided design. An H_0 -only corpus remains usable, but it supplies neither human dimensional judgments nor human explanations.

The learned object is an executable evaluation program

$$E_\theta = (A, K, V, T, \Pi, G), \quad (s_\theta, \tau_\theta) = E_\theta(x, y), \quad (2)$$

comprising artifact adapters A , dimension schema and rubrics K , evaluator modules V , tools T , evidence routing Π , and fusion G . The score s_θ is accompanied by an execution trace τ_θ . The base model is fixed outside θ . Self-evolution means automated changes to this evaluation program under external human feedback, not self-definition of the human standard. At deployment, the evaluator receives one candidate and its permitted evidence; opponents, human labels, generator identity, split labels, and mining strata are excluded. We seek agreement with unseen human preferences at an explicitly reported evaluation cost, rather than a single judge shared across all domains.

4 DISAGREEMENT-BASED PREFERENCE ACQUISITION

Preference acquisition allocates fewer annotations to comparisons on which existing evaluation signals consistently agree. It first makes these signals comparable, then identifies disagreements for human assessment while retaining random samples to check shared blind spots.

4.1 ALIGNING EVALUATION SIGNALS

For each domain, a versioned inventory links its quality dimensions to human protocols, rubrics, metrics, learned scorers, model judges, and analysis tools. Each source records its input requirements, provenance, limitations, and execution status. Only executable sources with compatible inputs supply signals; a coverage audit identifies dimensions lacking distinct sources. The comparison unit is an *output signal*, not an evaluator name. We freeze a many-to-many mapping M_{jq} from signal j to dimension q using its definition and measurement procedure before reading preference labels. A mixed score may connect to several dimensions, but remains one shared signal. Signals sharing a model, training labels, or underlying metric are assigned source groups so that duplicated prompts do not become independent votes. Existing executable metrics may enter the candidate tool library. Metrics reserved for final independent assessment are frozen separately and never enter tools, routing, critique, or reward.

4.2 IDENTIFYING DISAGREEMENTS

Candidate outputs are grouped by their original task input and deduplicated before comparison. For a pair $r = (x, y_a, y_b)$, signal j supplies a direction $u_j(r) \in \{-1, 0, +1, \perp\}$ from its own score difference and declared tie margin. Zero denotes a tie and $\perp$ missing or abstaining output; raw scores from different signals are never subtracted. Repeated-run variation and missingness are recorded separately from disagreement.

The miner distinguishes three cases. *Within-dimension disagreement* occurs when signals from distinct source groups favor opposite candidates on the same criterion. *Overall-dimension disagreement* occurs when an overall signal opposes a dimensional signal. *Cross-dimensional trade-offs* occur when dimension-level preferences favor different candidates; these are not contradictory measurements of one criterion. Dimension preferences aggregate at most one consensus direction per source group, with internally opposed groups abstaining. The implemented miner ranks pairs by

these structural conflicts, rather than treating disagreement as a known evaluator error. The priority score sums within-dimension and overall-dimension conflicts, gives the former twice the weight, and adds an indicator for cross-dimensional trade-offs. Appendix B specifies the ranking rule and separates this implementation from human-reliability calibration.

The acquisition design combines targeted comparisons with a random anchor selected before signal-based ranking. Human judgments on the anchor assess signal reliability and shared blind spots before the targeted policy is frozen. Anchor annotations count toward the same human budget as targeted annotations. Only evolution data are disagreement-sampled; final evaluation uses independently specified random or stratified sampling. The test of acquisition is subsequent evaluator improvement at equal annotation cost, not the number of disagreements collected.

4.3 FROM COMPARISONS TO EVOLUTION FEEDBACK

Existing preferences can be reused when the task, readable candidate outputs, and judged quality construct match. Derived comparisons retain their ranking, rating, or error-annotation provenance. When matching preferences are absent, the task and annotation protocol are fixed before collecting blinded judgments on existing or freshly generated candidates. The two settings retain separate provenance, agreement, ties, and abstentions (Appendix B).

The adaptation evidence table $Z(r, q)$ links artifact provenance, available human judgments, source signals, evaluator decisions, tool observations, and execution cost. Both construction settings supply this interface to ERA, while retaining their label origins and annotation protocols. Overall and dimensional preferences identify errors to investigate; localized diagnostics and human rationales are included when available. Missing feedback levels remain absent, and model-generated explanations are not human H_1 or H_2 .

5 ERA: EVIDENCE-GUIDED EVALUATOR EVOLUTION

ERA evolves an evaluation procedure by changing the program that implements it. Its central distinction is between accepting a candidate and continuing its direction: a failed candidate may still reveal how to make the intended improvement work.

5.1 EXECUTING AND MEASURING AN EVALUATOR

Each invocation evaluates one artifact through the program in Equation 2. Adapters expose permitted representations, such as source text, rendered views, or spatial structure. Evaluator modules apply their criteria and use routing rules to obtain tool observations, then return findings with supporting evidence and uncertainty. An available tool affects judgment only if its output reaches a module that uses it (Figure 3). The default semantic synthesis groups these findings by their consequences for quality and assigns an overall score on its own anchored scale. It receives specialist findings, reasons, and scopes, but not their structured scores, ratings, or score scales. Numerical fusion is an explicit alternative (Appendix E).

Full-evaluator development and terminal validation request $R \geq 3$ executions per artifact under a fixed protocol. Artifact scores are aggregated before comparison:

$$\bar{s}_\theta(x, y) = \text{median}_{r=1, \dots, R} s_\theta^{(r)}(x, y), \quad z_i = \mathbb{I}[\bar{s}_\theta(x_i, y_i^+) > \bar{s}_\theta(x_i, y_i^-)]. \quad (3)$$

Here y_i^+ denotes the human-preferred artifact, which is never identified to the scorer. Development agreement is the mean of z_i over the fixed cohort, with predicted ties counted as incorrect on strict-preference pairs. Missing executions retain their denominators and are reported through coverage and missing-outcome bounds; human ties have a separate prespecified policy (Appendix F).

5.2 TURNING AN EVALUATION ERROR INTO A CANDIDATE

A direction specifies an intended correction to the evaluation procedure; its hypothesis explains why that correction should improve a judgment. For example, detecting an important omission may require locating the missing source content, assessing its importance, and making the final score reflect that assessment.

The agent investigates errors using artifacts, human feedback, and traces from a proposal subset $\mathcal{D}_{\text{prop}}$. It can inspect component instructions and failed attempts to decide which changes are needed. In one tool-driven session, it can request additional $\mathcal{D}_{\text{prop}}$ measurements or component probes before submitting actual edits, a short intent, and references to inspected training cases. Expected orderings and intermediate observations can also be declared before measurement. A sealed training subset $\mathcal{D}_{\text{dev}}$, disjoint from $\mathcal{D}_{\text{prop}}$ by original input and artifact identity, returns aggregate development measurements without exposing its cases or labels.

A candidate is a complete evaluator version implementing a direction, not an individual edit. Edits can change any of (A, K, V, T, Π, G) while keeping the base model fixed. Connected changes, such as a tool, its route, and its consuming module, are assembled against one parent and evaluated as one complete candidate. Its outcome is not credited separately to every edit. The direction is the unit of continued exploration; the complete program is the unit of selection.

5.3 SEPARATING ACCEPTANCE FROM CONTINUATION

Search retains a best evaluator B and a working evaluator W for the active direction. A new direction starts from B ; subsequent attempts start from an eligible W . Depth-first continuation gives the active direction the next editing opportunity, using its diagnostics to address the remaining implementation problem until it is paused or search ends. Every valid, distinct candidate receives the common development measurement, irrespective of whether local diagnostics were favorable.

Candidate acceptance. A candidate replaces B only if development agreement strictly improves, execution is complete and comparable, coverage does not decrease, and the prespecified stratum-regression guard holds. These conditions select a development best; they are not a significance test.

Direction continuation. A measured candidate can become W without replacing B . Besides accepted candidates, complete and comparable candidates with no observed damage remain eligible. Damage is checked on local cases and in aggregate against both the parent and best evaluator. Retaining an unaccepted candidate with observed damage additionally requires new measured progress and no damage to declared protection cases. Progress means either a first, previously uncredited repair of an error present at the direction’s start or a previously uncredited trace-field change matching a predeclared exact value. Repairing damage introduced by an earlier attempt cannot repeatedly earn this credit. If inheritance is not justified, W returns to B , while the failed attempt remains available for diagnosis. Thus rejecting an implementation need not erase what search learned from it or end work on its direction. Algorithm 1 and Appendix D detail the loop and its checks.

5.4 INTERPRETING FAILURE AND ENDING SEARCH

Falsification guides the interpretation of a failed attempt: did the proposed behavior occur, and did it affect the judgment as hypothesized? In the omission example, finding missing content without using its importance in the score exposes an implementation obstacle without establishing that the direction is unhelpful. The next attempt can address that obstacle while preserving the best evaluator. A failed candidate does not by itself refute its direction’s hypothesis. Refutation requires diagnostic results that contradict the hypothesis within its stated scope after relevant implementation alternatives have been examined. The current controller implements measurement, inheritance, and continuation or pausing; it does not automatically adjudicate semantic refutation. Search is bounded by declared rounds and execution limits, and resource exhaustion leaves an unresolved direction rather than a refuted one.

On normal completion, eligible, previously accepted candidates are shortlisted and ordered using training measurements alone. Repeated access makes $\mathcal{D}_{\text{dev}}$ adaptive development data, not an independent test. Terminal validation $\mathcal{D}_{\text{val}}$ assesses this frozen shortlist without reranking it or feeding further edits; with no eligible improvement, search retains the seed. Independent Test/OOD assessment follows method and analysis freeze (Appendix C).

6 EXPERIMENTAL DESIGN

Core comparisons, ablations, and analysis address distinct evaluation questions. Domains enter confirmatory experiments after a readiness audit; models, data, human protocols, budgets, and primary contrasts are frozen before final evaluation. Every result is indexed by domain, data-construction setting, and ID/OOD source, with its provenance and annotation quality. Reuse tests improvement with existing labels; collection tests applicability when labels must be acquired. Neither estimates the causal effect of label origin.

6.1 CORE RESULTS

C1: evaluator alignment and selection. Frozen ERA evaluators are compared with the best single metric, a metric ensemble, a strong static judge, and a static tool-augmented judge. Prompt-optimization and program-search controls use declared versions and comparable edit interfaces. Fixed-structure evaluator fine-tuning and score calibration test parameter adaptation as alternatives (Appendix L). Matched comparisons share the base model, human supervision, inputs, candidate pools, and scoring protocol. Search cost and deployment cost are reported separately.

For each test input, all evaluators receive the same candidate group, with a default design of $N = 4$. Complete human rankings support global pairwise agreement, top-region agreement involving the human top-two candidates, Best-of- N selection, and selection regret. When only ranks are available, normalized regret is $(\text{rank}(\hat{y}) - 1)/(N - 1)$ for the selected output $\hat{y}$. A pair-only dataset supports pairwise evaluation, not an N -way selection claim. Core H_1 , worst-stratum performance, coverage, variance, and cost accompany overall agreement.

ID and OOD results are reported separately. OOD changes dataset or source distribution while preserving the evaluation standard. A difficult subset alone is a difficulty test, not an OOD source. Final sampling is independent of mining conflicts. All splitting and statistical resampling preserve original input/source groups. Paired effects use grouped 95% intervals. Existing annotations are reused only when their construct and candidate groups match the endpoint.

C2: critic-guided refinement. For each unseen input, a frozen generator produces one fresh output y_0 . The seed critic and ERA critic each inspect that same output and permitted evidence, then return the same critique format. A common generator performs one refinement step under identical prompts, decoding, tool access, and token budgets. Refinement validation fixes these settings before ID/OOD generation. New blinded human rankings assess y_0 , Seed-refined, and ERA-refined. The primary contrast is ERA-refined versus Seed-refined, accompanied by improvement and regression rates against y_0 , core dimensional judgments, and reserved external metrics. Labels on pre-existing outputs cannot assess these new outputs.

C3: reward-guided training. The planned study compares Base, Seed-GRPO, and ERA-GRPO from one checkpoint, using LoRA and group relative policy optimization (Hu et al., 2022; Shao et al., 2024). Training arms match prompts, optimizer settings, updates, policy compute, and reward-call budgets; scoring cost is reported separately. Rewards are frozen before training. Rollouts are sampled online by each arm’s own evolving policy and are not shared. New blinded H_0/H_1 and reserved metrics assess fresh ID/OOD outputs. Selection and reward-hacking safeguards appear in Appendix H. C3 remains a design-only study here; no parameter training or training benefit is reported.

6.2 ABLATION STUDIES

The ablations separate annotation choices from evaluator and search changes (Table 5). Mining strategies are compared at equal human budgets $N_H \in \{20, 50, 100, 200\}$, measured in original-input annotation units. Random and disagreement sampling use the same pool, feedback granularity, initial program, and evolution budget; reliability anchors are charged to every arm. Sampling rules are frozen before non-anchor labels are revealed, including in public-corpus replay. The endpoint is post-evolution alignment on an independently sampled cohort, not errors found during acquisition. Corpus construction alone cannot establish sampling efficiency. The value of H_1/H_2 requires a separate same-pair granularity comparison reporting both unit counts and annotation time. Com-




Figure 2: **Simulated task coverage and evaluator comparisons; not experimental results.** Panel (a) summarizes 22 tasks and 2,880 input groups, balanced to 720 per family with four candidates per input. A fixed hash selects this subset without reading outcomes; the appendix reports the full synthetic cohorts. Sector angles encode input counts; task codes follow Table 1. Panel (b) plots assigned cost units, not measured runtime, against OOD agreement with synthetic human rankings. Small marks denote 110 domain–method combinations; large marks equally average 11 domains per feedback setting. Shapes identify methods, and filled/open marks denote illustrative reused/new preferences, not label reliability. Panel (c) reports agreement from pointwise scores and 95% input-group bootstrap intervals on a common 0–100% scale. Baseline bars follow the displayed order in each family’s hue; ERA is lake green. Cost includes all selected inputs; agreement uses method-complete inputs. Marginal intervals do not establish paired differences.

ponent controls preserve unaffected interfaces and report invalid dependencies rather than reverting the whole program to its seed.

6.3 ANALYSIS AND STUDIES

Trajectory analysis records directions, implementation attempts, resource use, and evidence for scoped refutation. Component effects and remaining human-agreement gaps motivate tests for specification, observation, or headroom limits; strata within a domain may differ. Feedback analysis relates held-out changes to annotation units, time, and granularity. Adapter analysis compares direct judgments with labels derived from rankings, scores, or errors on the same subset. Repeated scoring and independent search runs remain distinct sources of variation.

7 QUANTITATIVE EVALUATION: SIMULATED DISPLAYS

All numerical displays in this section use *simulated data*, not measured evaluator performance or collected human judgments. They instantiate the reporting protocol (Appendix J); their assigned outcomes provide no evidence for either research hypothesis.

Alignment across evaluation standards. Average agreement can conceal opposing effects across tasks. Figure 2c retains all 22 domains in four families, comparing five evaluator fixtures on OOD inputs. Panel (a) reports their balanced input composition, while panel (b) relates agreement to

assigned evaluation cost without pooling the two feedback settings. Task order follows the taxonomy, not performance. The two metric-only references remain in Table 10; Figure 4 reports paired ERA-minus-seed effects on the full ID/OOD cohorts.

What changes in the evaluator, and what feedback supports it? Figure 8a reports six conditional full-minus-fixed component contrasts in eight diagnostic tasks, not additive contributions. Panel (b) varies cumulative feedback detail on the same input groups. Granularity comparisons must also account for annotation time and unequal label availability; they do not test acquisition efficiency.

Feedback utility, acquisition, and resources. Alignment does not establish useful critique. Figure 9a–b retains C2 quality distributions and unresolved judgments for the eight diagnostic tasks, on inputs distinct from C1. Both feedback settings require fresh blinded endpoint judgments for C2. Figure 5a–b isolates acquisition policy at equal human budgets, including a common random anchor. Figure 6a–c keeps agreement, abstention, and consensus denominators distinct; label origin does not establish reliability. Selection profiles (Figure 7) and deployment costs (Figure 10) retain separate denominators. C3 remains entirely unmeasured.

8 DISCUSSION

ERA makes the program’s approximation of human evaluation standards editable through feedback. Explicit criteria and evidence paths permit inspection, but whether they faithfully explain judgments requires intervention tests. Comparison with parameter adaptation asks whether evolving this structure is useful; not all disagreement is a structural failure.

8.1 LIMITATIONS AND FUTURE WORK

First, simulated displays support no effectiveness claim and must be replaced with frozen empirical cohorts. Second, dimensional judgments require additional annotation and construct validation. Third, shared signal bias and interacting components require dependence audits and controlled attribution. Finally, pointwise scoring assumes a useful within-domain ordering, while repeated development can overfit aggregate feedback. Independent ID/OOD and generation-level studies must test evaluation gains and downstream utility; transfer across standards remains a separate question.

9 CONCLUSION

Human preferences are difficult to express as executable evaluation procedures. We study evaluation-agent self-evolution through structured feedback and hypothesis-guided program changes, with a fixed base model. Reused and newly collected preferences support the same problem. The study design compares program evolution with parameter adaptation and separates preference alignment from usefulness in guiding generation. Simulated displays do not establish these benefits; the question is whether an evolved approximation of human standards remains useful on unseen outputs.

AI USE STATEMENT

Generative AI assisted manuscript drafting, literature organization, and the preparation of editable LaTeX conceptual diagrams and plotting code. GPT Image 2 generated the photograph in Figure 1a through scene generation and a controlled edit. The quantitative displays use explicitly simulated data, generated offline for reviewing the reporting design. They are not observed cases, human annotations, model executions, or performance measurements. No empirical results or parameter-training runs were produced for this draft. Bibliographic identities were checked against primary metadata sources. Author review, responsibility for the final content, and the final disclosure of research-stage AI use remain to be completed before submission.

ETHICS STATEMENT

Human preferences can encode annotator and model biases rather than a universal definition of quality. Domain-specific criteria and uncertainty should remain visible, especially for automated research judgments. Any new annotation study must document consent, compensation, privacy protection, and the applicable institutional review determination before data collection. Artifact licenses and source permissions must be checked independently of technical accessibility. Generated code and evidence tools require restricted execution and controlled external access. Reward-guided optimization requires checks for exploitation of the evaluator and degradation outside its favored criteria. No private artifacts, identifiable annotation records, or experimental traces are included in this working manuscript. Study-specific approvals and release decisions are pending.

REPRODUCIBILITY STATEMENT

Sections 3–6 define semantic mappings, feedback, the evaluator program, evolution, and controlled evaluation. Appendices B–F specify annotation, access boundaries, the algorithm, and measurement. Appendices G–K specify ablations, planned training, reporting outputs, simulated fixtures, and the experiment manifest. Appendix J identifies the offline generator and provenance for the numerical displays. The current manuscript contains no empirical results to reproduce. Before submission, the authors must supply a versioned, appropriately anonymized implementation and permitted data-processing and analysis artifacts with completed manifests. Their availability is not asserted here.

REFERENCES

- Lakshya A Agrawal, Shangyin Tan, Dilara Soylu, Noah Ziem, Rishi Khare, Krista Opsahl-Ong, Arnav Singhvi, Herumb Shandilya, Michael J Ryan, Meng Jiang, Christopher Potts, Koushik Sen, Alex Dimakis, Ion Stoica, Dan Klein, Matei Zaharia, and Omar Khattab. GEPA: Reflective Prompt Evolution Can Outperform Reinforcement Learning. In *International Conference on Learning Representations*, volume 2026, pp. 8479–8565, 2026. URL https://proceedings.iclr.cc/paper_files/paper/2026/hash/0e9e708b6f48e14fd0ac29e167413f76-Abstract-Conference.html.
- Xin-Sheng Chen, Jiayu Zhu, Pei lin Li, Hanzheng Wang, Shuojin Yang, and Meng-Hao Guo. PresentBench: A Fine-Grained Rubric-Based Benchmark for Slide Generation. *arXiv preprint arXiv:2603.07244*, 2026. URL <https://arxiv.org/abs/2603.07244>.
- Alexander R. Fabbri, Wojciech Kryściński, Bryan McCann, Caiming Xiong, Richard Socher, and Dragomir Radev. SummEval: Re-evaluating Summarization Evaluation. *Transactions of the Association for Computational Linguistics*, 9:391–409, 2021. doi: 10.1162/tacl_a_00373.
- Dhruba Ghosh, Hannaneh Hajishirzi, and Ludwig Schmidt. GenEval: An object-focused framework for evaluating text-to-image alignment. In *Advances in Neural Information Processing Systems*, volume 36, pp. 52132–52152, 2023. doi: 10.52202/075280-2270.
- Edward J. Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yanzhi Li, Shean Wang, Lu Wang, and Weizhu Chen. LoRA: Low-rank adaptation of large language models. In *International Con-*

- ference on Learning Representations, 2022. URL <https://openreview.net/forum?id=nZeVKeeFYf9>.
- Shengran Hu, Cong Lu, and Jeff Clune. Automated Design of Agentic Systems. In *International Conference on Learning Representations*, volume 2025, pp. 21344–21377, 2025. URL https://proceedings.iclr.cc/paper_files/paper/2025/hash/36b7acf6f6010652b3f2a433774a66fe-Abstract-Conference.html.
- Zisu Huang, Jingwen Xu, Yifan Yang, Ziyang Gong, Qihao Yang, Muzhao Tian, Xiaohua Wang, Changze Lv, Xuemei Gao, Qi Dai, Bei Liu, Kai Qiu, Xue Yang, Dongdong Chen, Xiaoqing Zheng, and Chong Luo. From Raw Experience to Skill Consumption: A Systematic Study of Model-Generated Agent Skills. *arXiv preprint arXiv:2605.23899*, 2026. URL <https://arxiv.org/abs/2605.23899>.
- Omar Khattab, Arnav Singhvi, Paridhi Maheshwari, Zhiyuan Zhang, Keshav Santhanam, Sri Vardhamanan, Saiful Haq, Ashutosh Sharma, Thomas T. Joshi, Hanna Moazam, Heather Miller, Matei Zaharia, and Christopher Potts. DSPy: Compiling Declarative Language Model Calls into Self-Improving Pipelines. In *International Conference on Learning Representations*, volume 2024, pp. 54928–54958, 2024. URL https://proceedings.iclr.cc/paper_files/paper/2024/file/f1cf02ce09757f57c3b93c0db83181e0-Paper-Conference.pdf.
- Seungone Kim, Jamin Shin, Yejin Cho, Joel Jang, Shayne Longpre, Hwaran Lee, Sangdoo Yun, Seongjin Shin, Sungdong Kim, James Thorne, and Minjoon Seo. Prometheus: Inducing Fine-grained Evaluation Capability in Language Models. In *International Conference on Learning Representations*, volume 2024, pp. 29927–29962, 2024a. URL https://proceedings.iclr.cc/paper_files/paper/2024/file/803485352e61e3ebf41221e4776c9fd4-Paper-Conference.pdf.
- Seungone Kim, Juyoung Suk, Shayne Longpre, Bill Yuchen Lin, Jamin Shin, Sean Welleck, Graham Neubig, Moontae Lee, Kyungjae Lee, and Minjoon Seo. Prometheus 2: An Open Source Language Model Specialized in Evaluating Other Language Models. In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*, pp. 4334–4353, 2024b. doi: 10.18653/v1/2024.emnlp-main.248.
- Yuval Kirstain, Adam Polyak, Uriel Singer, Shahbuland Matiana, Joe Penna, and Omer Levy. Pick-a-Pic: An Open Dataset of User Preferences for Text-to-Image Generation. In *Advances in Neural Information Processing Systems*, volume 36, pp. 36652–36663, 2023. doi: 10.52202/075280-1594.
- Nathan Lambert, Valentina Pyatkin, Jacob Morrison, LJ Miranda, Bill Yuchen Lin, Khyathi Chandu, Nouha Dziri, Sachin Kumar, Tom Zick, Yejin Choi, Noah A. Smith, and Hannaneh Hajishirzi. RewardBench: Evaluating Reward Models for Language Modeling. In *Findings of the Association for Computational Linguistics: NAACL 2025*, pp. 1755–1797. Association for Computational Linguistics, 2025. doi: 10.18653/v1/2025.findings-naacl.96. URL <https://aclanthology.org/2025.findings-naacl.96/>.
- Yang Liu, Dan Iter, Yichong Xu, Shuohang Wang, Ruochen Xu, and Chenguang Zhu. G-Eval: NLG Evaluation using GPT-4 with Better Human Alignment. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, pp. 2511–2522, 2023. doi: 10.18653/v1/2023.emnlp-main.153.
- Yuxuan Liu, Tianchi Yang, Shaohan Huang, Zihan Zhang, Haizhen Huang, Furu Wei, Weiwei Deng, Feng Sun, and Qi Zhang. Calibrating LLM-based evaluator. In *Proceedings of the 2024 Joint International Conference on Computational Linguistics, Language Resources and Evaluation (LREC-COLING 2024)*, pp. 2638–2656. ELRA and ICCL, 2024. doi: 10.63317/3posqwchnb6g. URL <https://aclanthology.org/2024.lrec-main.237/>.
- Changze Lv, Jie Zhou, Wentao Zhao, Jingwen Xu, Shihan Dou, Zisu Huang, Muzhao Tian, Xiaohua Wang, Zhengkang Guo, Yang Liu, Pluto Zhou, Tao Gui, Le Tian, Xiao Zhou, Xiaoqing Zheng, Xuanjing Huang, and Jie Zhou. Learning Query-Specific Rubrics from Human Preferences for DeepResearch Report Generation. *arXiv preprint arXiv:2602.03619*, 2026. URL <https://arxiv.org/abs/2602.03619>.

- Aman Madaan, Niket Tandon, Prakhar Gupta, Skyler Hallinan, Luyu Gao, Sarah Wiegrefe, Uri Alon, Nouha Dziri, Shrimai Prabhumoye, Yiming Yang, Shashank Gupta, Bodhisattwa Prasad Majumder, Katherine Hermann, Sean Welleck, Amir Yazdanbakhsh, and Peter Clark. Self-Refine: Iterative Refinement with Self-Feedback. In *Advances in Neural Information Processing Systems*, volume 36, pp. 46534–46594, 2023. doi: 10.52202/075280-2019.
- William Muldrew, Peter Hayes, Mingtian Zhang, and David Barber. Active preference learning for large language models. In *Proceedings of the 41st International Conference on Machine Learning*, volume 235 of *Proceedings of Machine Learning Research*, pp. 36577–36590. PMLR, 2024. URL <https://proceedings.mlr.press/v235/muldrew24a.html>.
- Long Ouyang, Jeffrey Wu, Xu Jiang, Diogo Almeida, Carroll Wainwright, Pamela Mishkin, Chong Zhang, Sandhini Agarwal, Katarina Slama, Alex Ray, John Schulman, Jacob Hilton, Fraser Kelton, Luke Miller, Maddie Simens, Amanda Askell, Peter Welinder, Paul F Christiano, Jan Leike, and Ryan Lowe. Training language models to follow instructions with human feedback. In *Advances in Neural Information Processing Systems*, volume 35, pp. 27730–27744, 2022. doi: 10.52202/068431-2011.
- Reid Pryzant, Dan Iter, Jerry Li, Yin Lee, Chenguang Zhu, and Michael Zeng. Automatic Prompt Optimization with “Gradient Descent” and Beam Search. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, pp. 7957–7968, 2023. doi: 10.18653/v1/2023.emnlp-main.494.
- Qi Qian, Chengsong Huang, Jingwen Xu, Changze Lv, Muling Wu, Wenhao Liu, Xiaohua Wang, Zhenghua Wang, Zisu Huang, Muzhao Tian, Jianhan Xu, Kun Hu, He-Da Wang, Yao Hu, Xu-anjing Huang, and Xiaoqing Zheng. Benchmark²: Systematic Evaluation of LLM Benchmarks. *arXiv preprint arXiv:2601.03986*, 2026. URL <https://arxiv.org/abs/2601.03986>.
- Rafael Rafailov, Archit Sharma, Eric Mitchell, Christopher D Manning, Stefano Ermon, and Chelsea Finn. Direct Preference Optimization: Your Language Model is Secretly a Reward Model. In *Advances in Neural Information Processing Systems*, volume 36, pp. 53728–53741, 2023. doi: 10.52202/075280-2338.
- H. S. Seung, M. Opper, and H. Sompolinsky. Query by committee. In *Proceedings of the Fifth Annual Workshop on Computational Learning Theory*, pp. 287–294. ACM, 1992. doi: 10.1145/130385.130417. URL <https://doi.org/10.1145/130385.130417>.
- Shreya Shankar, J.D. Zamfirescu-Pereira, Bjoern Hartmann, Aditya Parameswaran, and Ian Arawjo. Who Validates the Validators? Aligning LLM-Assisted Evaluation of LLM Outputs with Human Preferences. In *Proceedings of the 37th Annual ACM Symposium on User Interface Software and Technology*, pp. 1–14, 2024. doi: 10.1145/3654777.3676450.
- Zhihong Shao, Peiyi Wang, Qihao Zhu, Runxin Xu, Junxiao Song, Xiao Bi, Haowei Zhang, Mingchuan Zhang, Y. K. Li, Y. Wu, and Daya Guo. DeepSeekMath: Pushing the limits of mathematical reasoning in open language models. *arXiv preprint arXiv:2402.03300*, 2024. URL <https://arxiv.org/abs/2402.03300>.
- Noah Shinn, Federico Cassano, Ashwin Gopinath, Karthik Narasimhan, and Shunyu Yao. Reflexion: language agents with verbal reinforcement learning. In *Advances in Neural Information Processing Systems*, volume 36, pp. 8634–8652, 2023. doi: 10.52202/075280-0377.
- Chenglei Si, Yanzhe Zhang, Ryan Li, Zhengyuan Yang, Ruibo Liu, and Diyi Yang. Design2Code: Benchmarking Multimodal Code Generation for Automated Front-End Engineering. In *Proceedings of the 2025 Conference of the Nations of the Americas Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 1: Long Papers)*, pp. 3956–3974, 2025. doi: 10.18653/v1/2025.naacl-long.199.
- Nisan Stiennon, Long Ouyang, Jeffrey Wu, Daniel Ziegler, Ryan Lowe, Chelsea Voss, Alec Radford, Dario Amodei, and Paul F Christiano. Learning to summarize with human feedback. In *Advances in Neural Information Processing Systems*, volume 33, pp. 3008–3021, 2020. URL https://proceedings.neurips.cc/paper_files/paper/2020/file/1f89885d556929e98d3ef9b86448f951-Paper.pdf.

- Sijun Tan, Siyuan Zhuang, Kyle Montgomery, William Y. Tang, Alejandro Cuadron, Chenguang Wang, Raluca Ada Popa, and Ion Stoica. JudgeBench: A Benchmark for Evaluating LLM-based Judges. In *International Conference on Learning Representations*, volume 2025, pp. 63277–63303, 2025. URL https://proceedings.iclr.cc/paper_files/paper/2025/file/9e720f91114c49cf9d640d821da3-Paper-Conference.pdf.
- Tianlu Wang, Ilia Kulikov, Olga Golovneva, Ping Yu, Weizhe Yuan, Jane Dwivedi-Yu, Richard Yuanzhe Pang, Maryam Fazel-Zarandi, Jason Weston, and Xian Li. Self-Taught Evaluators. *arXiv preprint arXiv:2408.02666*, 2024. URL <https://arxiv.org/abs/2408.02666>.
- Xiaohua Wang, Muzhao Tian, Yuqi Zeng, Zisu Huang, Jiakang Yuan, Bowen Chen, Jingwen Xu, Mingbo Zhou, Wenhao Liu, Muling Wu, Zhengkang Guo, Qi Qian, Yifei Wang, Feiran Zhang, Ruicheng Yin, Shihan Dou, Changze Lv, Tao Chen, Kaitao Song, Xu Tan, Tao Gui, Xiaoqing Zheng, and Xuanjing Huang. Reward Hacking in the Era of Large Models: Mechanisms, Emergent Misalignment, Challenges. *arXiv preprint arXiv:2604.13602*, 2026. URL <https://arxiv.org/abs/2604.13602>.
- Jiazhen Xu, Xiao Liu, Yuchen Wu, Yuxuan Tong, Qinkai Li, Ming Ding, Jie Tang, and Yuxiao Dong. ImageReward: Learning and Evaluating Human Preferences for Text-to-Image Generation. In *Advances in Neural Information Processing Systems*, volume 36, pp. 15903–15935, 2023. doi: 10.52202/075280-0700.
- Chengrun Yang, Xuezhi Wang, Yifeng Lu, Hanxiao Liu, Quoc V. Le, Denny Zhou, and Xinyun Chen. Large Language Models as Optimizers. In *International Conference on Learning Representations*, volume 2024, pp. 12028–12068, 2024. URL https://proceedings.iclr.cc/paper_files/paper/2024/hash/3339f19c5fcee3ad74502947a32be9e6-Abstract-Conference.html.
- Shunyu Yao, Jeffrey Zhao, Dian Yu, Nan Du, Izhak Shafran, Karthik Narasimhan, and Yuan Cao. ReAct: Synergizing Reasoning and Acting in Language Models. In *International Conference on Learning Representations*, 2023. URL <https://arxiv.org/abs/2210.03629>.
- Mert Yuksekogul, Federico Bianchi, Joseph Boen, Sheng Liu, Pan Lu, Zhi Huang, Carlos Guestrin, and James Zou. Optimizing generative AI by backpropagating language model feedback. *Nature*, 639(8055):609–616, 2025. doi: 10.1038/s41586-025-08661-4.
- Jenny Zhang, Shengran Hu, Cong Lu, Robert Lange, and Jeff Clune. Darwin Gödel Machine: Open-Ended Evolution of Self-Improving Agents. In *International Conference on Learning Representations*, volume 2026, pp. 104223–104294, 2026. URL https://proceedings.iclr.cc/paper_files/paper/2026/hash/aa5f5e6eb6f613ec412f1d948dfa21a5-Abstract-Conference.html.
- Jiayi Zhang, Jinyu Xiang, Zhaoyang Yu, Fengwei Teng, XiongHui Chen, Jiaqi Chen, Mingchen Zhuge, Xin Cheng, Sirui Hong, Jinlin Wang, Bingnan Zheng, Bang Liu, Yuyu Luo, and Chenglin Wu. AFlow: Automating Agentic Workflow Generation. In *International Conference on Learning Representations*, volume 2025, pp. 34040–34077, 2025. URL <https://openreview.net/forum?id=z5uVAKwmjf>.
- Tianyi Zhang, Varsha Kishore, Felix Wu, Kilian Q. Weinberger, and Yoav Artzi. BERTScore: Evaluating Text Generation with BERT. In *International Conference on Learning Representations*, 2020. URL <https://arxiv.org/abs/1904.09675>.
- Hao Zheng, Xinyan Guan, Hao Kong, Wenkai Zhang, Jia Zheng, Weixiang Zhou, Hongyu Lin, Yaojie Lu, Xianpei Han, and Le Sun. PPTAgent: Generating and Evaluating Presentations Beyond Text-to-Slides. In *Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing*, pp. 14402–14418, 2025. doi: 10.18653/v1/2025.emnlp-main.728.
- Lianmin Zheng, Wei-Lin Chiang, Ying Sheng, Siyuan Zhuang, Zhanghao Wu, Yonghao Zhuang, Zi Lin, Zhuohan Li, Dacheng Li, Eric Xing, Hao Zhang, Joseph Gonzalez, and Ion Stoica. Judging LLM-as-a-Judge with MT-Bench and Chatbot Arena. In *Advances in Neural Information Processing Systems*, volume 36, pp. 46595–46623, 2023. doi: 10.52202/075280-2020.

Mingchen Zhuge, Changsheng Zhao, Dylan R. Ashley, Wenyi Wang, Dmitrii Khizbullin, Yungang Xiong, Zechun Liu, Ernie Chang, Raghuraman Krishnamoorthi, Yuandong Tian, Yangyang Shi, Vikas Chandra, and Jürgen Schmidhuber. Agent-as-a-judge: Evaluate agents with agents. In *Proceedings of the 42nd International Conference on Machine Learning*, volume 267 of *Proceedings of Machine Learning Research*, pp. 80569–80611. PMLR, 2025. URL <https://proceedings.mlr.press/v267/zhuge25a.html>.

A DOMAIN PANORAMA AND EVALUATION STANDARDS

The research scope is organized by evaluation standard, not by file format. Table 1 retains all 22 candidate domains. Its dimensions are a starting taxonomy to be refined from domain evaluation knowledge before selecting datasets. Each instantiated domain must additionally fix priorities, decision rules, and an annotation protocol. Inclusion in this panorama does not establish data readiness, an implemented adapter, or empirical transfer. Figure 1a gives representative artifacts; the table makes their differing evaluation criteria explicit.

Table 1: Complete domain panorama. The four families organize the scope; every row requires its own standard and readiness audit. These are study candidates, not 22 completed evaluations.

ID	Domain	Core quality considerations
Structured visual artifacts		
TD-SV1	Brief-conditioned single-slide generation	Brief coverage, hierarchy, readability, composition, and visual consistency.
TD-SV2	Screenshot-conditioned webpage reconstruction	Reference content, spatial arrangement, typography, and visual fidelity.
TD-SV3	Data-to-visualization generation	Data fidelity, valid encodings, readability, and communicative usefulness.
TD-SV4	Specification-conditioned webpage generation	Requirement satisfaction, functional behavior, usability, and visual organization.
TD-SV5	Source-grounded slide-deck generation	Source faithfulness, narrative, coverage, slide readability, and deck consistency.
TD-SV6	Topic-conditioned slide-deck generation	Relevance, factual support, narrative completeness, audience fit, and presentation quality.
Image generation and understanding		
TD-IG1	Detailed image-caption generation	Grounding, coverage, relation accuracy, specificity, and absence of hallucination.
TD-IG2	Text-to-image generation	Prompt alignment, object and relation accuracy, visual coherence, and preference.
TD-IG3	Single-source instruction-based image editing	Instruction fulfillment, preservation of unrequested content, and image quality.
TD-IG4	LR-conditioned image super-resolution	Reference fidelity when available, perceptual detail, artifact avoidance, and plausible texture.
TD-IG5	Text-to-3D object generation	Prompt correspondence, geometry, appearance, and multi-view consistency.
TD-IG6	Image-to-3D object generation	Source-view fidelity, plausible unseen geometry, appearance, and multi-view consistency.
Text generation		
TD-TG1	Source-grounded document summarization	Faithfulness, relevance, coverage, coherence, and concision.
TD-TG2	RAG-supported long-form answer generation	Answer correctness, evidence grounding, attribution, completeness, and organization.
TD-TG3	Machine translation	Meaning preservation, omissions and additions, fluency, terminology, and style.
TD-TG4	Instruction-constraint-following response generation	Constraint satisfaction, substantive correctness, relevance, and response quality.
TD-TG5	Open-ended dialogue-response generation	Contextual appropriateness, helpfulness, consistency, safety, and conversational quality.
Automated research		
TD-AR1	Research-idea generation	Novelty relative to evidence, significance, feasibility, and testability.
TD-AR2	ML/NLP ablation-plan generation	Hypothesis identification, controlled contrasts, confound management, and feasibility.
TD-AR3	Table-to-insight report generation	Numerical faithfulness, appropriate inference, uncertainty, and relevance.

Table 1 continued

ID	Domain	Core quality considerations
TD-AR4	Related-work generation	Coverage, citation correctness, faithful comparison, and synthesis.
TD-AR5	Peer-review critique generation	Technical accuracy, evidence, significance assessment, fairness, and actionable criticism.

Representative resources include PPTAgent and PresentBench for presentations, Design2Code for screenshot-conditioned webpages, SummEval for summarization, and image-preference and alignment resources (Zheng et al., 2025; Chen et al., 2026; Si et al., 2025; Fabbri et al., 2021; Kirstain et al., 2023; Xu et al., 2023; Ghosh et al., 2023). Their outputs, references, metrics, and annotations have different evidential roles; a resource is not automatically a source of the required human preference. A benchmark collection spanning several standards remains a dataset collection, not an additional domain. Research execution with a complete executable postcondition is excluded, while ideation, study design, interpretation, writing, and critique remain open-ended evaluation problems.

Dataset roles. Dev supports acquisition and evaluator development, ID tests unseen inputs from the development source distribution, and OOD changes source distribution while preserving the standard. External datasets provide separately locked assessment only and do not determine semantic mappings or acquisition. Harder examples from the same source are a difficulty stratum, not a new OOD domain. A domain selected for confirmatory execution must pass the audit below; selection does not remove unselected domains from the research scope.

B FEEDBACK PROVENANCE, ANNOTATION, AND READINESS

The two construction settings in Table 2 share the same evaluator-evolution problem. They differ in the work needed to obtain matched human feedback. A domain may use both settings. Direct, derived, and materials-only provenance describe what a resource supplies; they are not additional domain classes or alternative algorithms.

Inputs, references, editable artifacts, and automatic metrics without matched human judgments are materials only. They may support candidate construction or independent assessment, but do not supply human supervision. When H_0 is absent, we preserve the original task and evaluation standard, separate source groups, recover readable candidates or generate them with frozen generators, and collect the required feedback. Preparation can precede annotation; final human-alignment claims cannot precede matched human assessment.

Annotation units and levels. One annotation unit is an original input and its anonymized candidate group, including permitted references. Pairs derived from that group remain dependent. The core tuple in Equation 1 is retained together with dataset and license, source grouping, generator and prompt version, collection time, annotator pool, agreement, and conversion provenance. Generator identity is recorded for audit but hidden from annotators and the deployed scorer.

The planned evolution protocol collects H_0 and core H_1 , with H_2 on conflict and failure subsets. H_3 is optional. Two independent annotators with adjudication may support development; final ID/OOD assessment targets three independent annotators and blinded complete rankings. Core dimensions and the tie policy are fixed in advance. A model or an author decision is not substituted for a missing independent final annotation. An H_0 -only instantiation must be reported as such; the framework’s richer interface does not establish that a run actually received H_1 or H_2 .

Agreement, abstention, and reannotation. New labels are not assumed to be more reliable than reused labels. Report the agreement statistic, its valid-response denominator, and the number of raters per group. Abstention uses all requested ratings as its denominator, including those excluded from valid-pair agreement. Group outcomes distinguish decisive preference, tie, no consensus, and insufficient votes. The simulated audit in Figure 6a–c illustrates these separate quantities. Individual disagreements are retained rather than filtered to match an old winner.

Table 2: Two routes to human feedback. Report each source’s protocol, agreement, abstention, and available supervision separately. An adapter changes representation, not the quality construct or the provenance of a judgment.

Setting	Label origin and protocol	Permitted supervision
Reused preferences	Original human pairs/rankings on readable same-input candidates; retain the source protocol and document any conversion from rankings, scores, or errors	H_0 only for overall judgments; H_1/H_2 only when actually annotated. Unknown source agreement or abstention remains unknown
New preferences	Freeze the task, visible evidence, dimensions, ties, and abstention before collecting independent blind judgments on existing or freshly generated candidates	Collect the levels required by the study; preserve individual votes and unresolved disagreements. Fresh C2/C3 outputs require new endpoint judgments

Annotators do not see model scores, method identity, or earlier winners. A changed rubric or visible evidence defines a changed protocol; resulting differences do not by themselves establish errors in the older labels.

Label-blind acquisition. Freeze domain knowledge, signal definitions, semantic mappings, and the calibration recipe before accessing preference labels. Select a random anchor, reveal only its authorized training judgments to estimate reliability, then freeze the targeted sampling policy before opening labels in the remaining pool. Charge anchor annotations to every strategy’s human budget. An existing labeled dataset must simulate the same access order. The final cohort is sampled independently of conflicts and never contributes to reliability calibration.

Implemented conflict ranking. The miner computes structural disagreement without reading the pair’s human preference. Let $I_{\text{within},q}(r)$ indicate opposing signals from distinct source groups on dimension q , $I_{\text{overall},q}(r)$ indicate opposition between an overall and a dimensional signal, and $I_{\text{tradeoff}}(r)$ indicate opposed dimension-level preferences. Its current priority rule is

$$P(r) = 2 \sum_q I_{\text{within},q}(r) + \sum_q I_{\text{overall},q}(r) + I_{\text{tradeoff}}(r). \quad (4)$$

The coefficients define a structural ranking, not estimated signal reliability or a probability of evaluation error. Overall–dimension comparisons retain whether their signals share a source group; they are not automatically independent evidence. Each source group contributes at most one direction to a dimension’s aggregate preference, with opposing within-group directions yielding abstention. Missing outputs and small-margin ties remain distinct from opposing preferences.

For repeated model signals, the acquisition margin is the median of same-repeat score differences between the two artifacts; variation and sign changes are retained. This differs from Equation 3, which first takes each artifact’s median final score and then compares them. A missing required signal repeat makes that acquisition signal unavailable. Deterministic metrics use declared ranges and orientations, not ranges fitted to the candidate pool or human labels. The random anchor is selected by original-input hashes before scoring, preserving candidate groups. Label-blind execution does not establish label-blind upstream cohort selection; historical data must also document how the candidate pool was chosen.

Human qualification and implementation boundary. The implemented miner supplies the structural queue and its source-dependence records. The full acquisition protocol additionally requires human-annotated anchors to qualify signals before targeted annotation; cross-domain human-reliability calibration and the equal-budget advantage remain unverified. Model-annotated pilot anchors measure agreement with that model, not human reliability. The annotation fraction, qualification thresholds, and targeted policy must be frozen for the human study, rather than inherited silently from a pilot. The evidence table Z is a logical join over feedback, artifact, signal, and execution records; absent human dimensions or rationales remain absent.

Source and tool status. Registry entries distinguish executable, planned, and blocked sources (Table 3). A publication or code link without compatible inputs, weights, or dependencies is not an

Table 3: Readiness criteria before a domain enters confirmatory evaluation. These are audit requirements, not completed coverage scores.

Audit object	Required record and exclusion condition
Evaluation standard	Open-endedness, dimensions and hierarchy, priorities, decision rule, and human protocol. Exclude tasks fully decided by an executable postcondition.
Knowledge provenance	Benchmarks, expert guidance, rubrics, and methods with versions, dates, input assumptions, known biases, and availability.
Source coverage	Dimension $\times$ source-type matrix: deterministic, learned metric, LLM/VLM judge, external analysis tool. Each core dimension needs executable evidence.
Source independence	For conflict mining, ordinarily two mechanisms per dimension. Record shared models, training labels, metrics, and mixed signals; multiple names are not independent votes.
Executable artifacts	Recoverable input/candidate groups, readable artifacts, references, permissions, and source fingerprints. Missing assets do not become scorer disagreement.
Conflict pool	Three semantic strata, availability, margins, repeat variance, anchor reliability, dependence caps, and uncovered strata. Noise alone is not a gap.
Feedback fidelity	Direct versus derived provenance, agreement, mapping to $H_0/H_1/H_2$, adapter checks on a common subset, and additional annotation needs.
Source isolation	Original-input deduplication, Dev/ID/OOD/External roles, contamination review, and a pre-search External Metric Reserve.

Table 4: Information available to evolution and assessment. Every invocation of the deployed scorer remains pointwise and label-blind, including on Train.

Partition	Information released	Permitted decision
Proposal Train	Authorized artifacts, feedback, and traces	Mechanism hypotheses and candidate evolution
Sealed Train	Prespecified aggregates only	Development promotion, continuation, and allocation
Terminal Val	Frozen-shortlist measurements	Terminal assessment; no reranking or further edits
Test/OOD	Locked until method and analysis freeze	Independent final reporting
Metric reserve	No access during evolution	Final assessment only; no tools, critique, reward, or selection

executed signal. Available tools begin as candidates; a route must explicitly activate them. Source independence is assessed at the signal level, including model, training-label, and metric dependencies. Model-only comparisons hold rubric, evidence, context, and reasoning budget fixed. Replacing both the scoring model and its evidence-collection agent instead compares full systems. Deterministic or fixed learned metrics are not duplicated for every judge model.

C DATA ACCESS AND INDEPENDENT ASSESSMENT

Table 4 separates proposal evidence, adaptive development measurements, and terminal assessment. The scorer always receives the same pointwise input contract, regardless of the cohort in which it is invoked. Only the authorized diagnostic process sees human labels and pair identities.

Training and validation. Proposal-Train evidence may inform evolution directions, candidates, and allocation. Sealed Train reveals prespecified aggregate agreement, coverage, and stratum guards without examples, labels, or per-case errors. Aggregates remain adaptive development evidence, not a final performance estimate. The terminal shortlist and its order are fixed from training evidence before Val access, and every shortlisted outcome is reported. Repeated Val calls can assess stability but do not create an independent test cohort. Interrupted searches do not obtain a new terminal assessment merely because they stopped.

External metrics. The External Metric Reserve is frozen before search. Reserved metrics and their equivalent wrappers are excluded from the executable tool library, router, diagnostician, critique, reward, and checkpoint-selection criteria. If an endpoint was used for optimization, it loses

independent-assessment status and must be replaced or explicitly reported as an optimization metric. This restriction does not forbid using other existing metrics as tools.

D ERA ALGORITHM AND ATTRIBUTION CONTRACT

Figure 3 separates the editable observation and judgment mechanisms. Algorithm 1 specifies their evolution and selection.




Figure 3: The six-component evaluation program. (a) Read one candidate and permitted task evidence. (b) Route tool observations. (c) Apply dimensional criteria. (d) Fuse sourced findings into a score and critique. The lower rail identifies the editable components. Human feedback supervises evolution, never the deployed scorer. Evidence and judgment remain separate; the summary and findings are a constructed example.

Algorithm 1 ERA: separate candidate acceptance from direction continuation

Require: Seed E_0 , proposal evidence Z_{prop} , sealed $\mathcal{D}_{\text{dev}}$, fixed search and execution limits

```

1: Measure  $E_0$ ; initialize best  $B \leftarrow E_0$ , working  $W \leftarrow E_0$ , and history
2: while search resources remain do
3:   Start a direction from  $B$ , or continue the active direction from  $W$ 
4:   Inspect proposal-Train artifacts and history; request diagnostics as needed
5:   Construct complete candidate  $C$  with its intent and inspected Train references
6:   if  $C$  is valid and distinct then
7:     Measure  $C$  on  $\mathcal{D}_{\text{dev}}$ ; retain local diagnostic outcomes
8:     if the candidate-acceptance conditions hold then
9:        $B \leftarrow C$ 
10:    end if
11:     $W \leftarrow C$  if measured inheritance is eligible; otherwise  $W \leftarrow B$ 
12:  end if
13:  Retain changes, outcomes, failed attempts, and resource use in history
14:  Continue the direction with its diagnostics, or pause it
15: end while
16: On normal completion, freeze and order eligible candidates using Train only
17: Assess a nonempty shortlist on terminal Val without reranking or further edits
18: Retain the Train-selected evaluator, or the seed if no eligible improvement exists

```

Investigation and candidate construction. The evolution agent reads original proposal-Train artifacts, installed component instructions, execution traces, and relevant history through a tool-driven editing session. It can request additional authorized proposal-Train measurements or fixed-input component probes. The submission contains actual edits, a short intent, and references to inspected training cases. Predicted orderings, contrast cases, and intermediate checks are optional declarations, not controller-inferred scientific predictions. The controller records the exact parent, installed program, and measured outcomes. Local diagnostic success is not an admission requirement for the common development measurement.

An elementary edit changes one component; a connected candidate can assemble a tool, its route, and its consumer against a fixed parent. Only the complete candidate is executed and measured, not its partially assembled versions. Changed tools must be connected to valid routes, although valid routing alone does not establish correct use of their outputs. Interface validity, fixed-model constraints, and execution limits remain mandatory; edit-character counts are recorded as guidance rather than a hard acceptance or stopping rule. Unchanged programs and verified historical replays are distinguished from fresh trials, and failed calls remain in the cost record.

Diagnostic observations. Fixed-input probes compare supported specialist or synthesis variants while holding their upstream inputs fixed. Such probes bind their observations to the exact tested candidate and inputs; changing a tool, route, or topology requires the corresponding complete execution. An optional exact intermediate check records whether a declared field changed to its expected value on at least one complete repeat. All repeat-level matches, contradictions, and unknowns remain visible. A semantic expectation is not automatically verified, and a changed recorded field is not proof of a correct causal explanation or improved overall judgment.

Acceptance and working-state inheritance. Candidate acceptance requires complete, comparable execution, a strict development-agreement gain over the best evaluator, nondecreasing coverage, and the prespecified worst-stratum guard. These are development selection rules rather than statistical significance claims. Working-state inheritance is a separate decision. Accepted candidates and candidates with no observed damage can be inherited when their required measurements are complete and comparable. For this decision, damage includes a local ordering regression, deterioration on a declared protection case, or more regressed than corrected sealed-Train comparisons against either the parent or the best evaluator. An unaccepted candidate with damage is eligible only when fresh measurement supplies new progress and the declared protection cases do not deteriorate.

Qualifying progress is either a first-time repair of an error present at the direction’s root, still wrong in the parent, or a newly observed, predeclared exact intermediate-field change. Previously credited repairs and checks are retained across attempts, so breaking and repairing an already corrected case cannot renew this allowance. A missing protected case is unknown, not preserved. When inheritance fails, the next parent returns to the best evaluator, but the failed candidate and its diagnostics remain available. Keeping an intermediate candidate never changes the recorded best score.

Failure interpretation and termination. The implemented controller records continuation, revision, pausing, and resource limits; it does not assign a semantic refutation automatically. A research claim of scoped refutation must identify the contradicted prediction and examine relevant alternatives such as missing observations, unused tool output, or incorrect downstream interpretation. One unsuccessful implementation, a nonsignificant difference, or budget exhaustion alone is insufficient. These evidence requirements govern interpretation of the search record, not an unimplemented automatic hypothesis test.

Terminal assessment and attribution. The terminal shortlist includes eligible, previously accepted, complete programs with positive development gain over the seed and is ordered using Train evidence alone. Its membership and order are frozen before Val access. Terminal Val does not rerank the shortlist or guide another proposal. With no eligible improvement, the seed is retained without an additional Val query. Fresh system confirmation and independent Test/OOD assessment remain separate from adaptive development.

A connected candidate receives one system-level outcome, not independent causal credit for each component. Component attribution requires paired interventions with the remaining interfaces fixed, such as tool on/off or evolved versus seed synthesis. If removing a dependency makes the program invalid, the comparison must report that limitation rather than substitute an unrelated whole-program control. Complete prompt templates, interaction traces, candidate changes, and execution settings belong in the reproducibility record rather than the main algorithm.

E SEMANTIC SYNTHESIS AND ANCHORED SCORING

The default fusion procedure forms an overall judgment from specialist findings; a declared numerical rule is an explicit alternative. Each specialist produces findings with identifiers, supporting

evidence, dimensions, and uncertainty, as well as its own rating and reason. The full specialist trace remains available for diagnosis, but the synthesizer’s default input omits the structured specialist score, rating, and score scale. It retains the findings, reasons, and scope of each specialist and applies its own overall scale. This input separation does not remove value judgments expressed in the reasons.

The synthesizer groups findings by their consequences for quality, including strengths, weaknesses, mixed effects, and uncertainty. Every finding reference must occur in exactly one group. This accounting rule makes omitted or multiply referenced findings detectable; it does not guarantee correct semantic grouping. Duplicate descriptions should not become independent penalties, and favorable style should not cancel a factual error when the domain standard prioritizes faithfulness.

Let $a_k < a_{k+1}$ be adjacent overall anchors and $\lambda \in [0, 1]$ locate quality between them. The resulting scalar is

$$s = (1 - \lambda)a_k + \lambda a_{k+1}. \quad (5)$$

The anchors have domain-specific verbal meanings; specialist scores with different meanings are not averaged. Reference completeness, semantic correctness, and calibration require separate checks. A synthetic arithmetic example uses $a_k = 50$, $a_{k+1} = 75$, and $\lambda = 0.4$, giving $s = 60$. These numbers illustrate interpolation, not a universal scale or an evaluator result.

F MEASUREMENT AND STATISTICAL ANALYSIS

Completeness and denominators. For a prespecified strict-preference cohort of N pairs, let $c_i = 1$ when both artifacts have all R valid complete scores. Let $m = \sum_i c_i$ and $k = \sum_{i:c_i=1} z_i$, with z_i from Equation 3. Report completed-pair accuracy k/m when $m > 0$, coverage m/N , and the missing-outcome interval

$$\left[\frac{k}{N}, \frac{k + N - m}{N} \right]. \quad (6)$$

This interval bounds full-cohort agreement under arbitrary missing outcomes; it is not a statistical confidence interval. A missing repeat is not replaced by the median of surviving calls. Completeness includes selected routes and tools, not only a final score. A valid observation that an artifact fails a quality check differs from an execution error.

Ranking and selection endpoints. For the default $N = 4$ design, human complete rankings support all strict pairs. Global agreement averages within input groups before aggregation across inputs. Top-region agreement uses prespecified strict comparisons involving at least one human top-two candidate. Best-of- N measures selection of a human top-ranked candidate; normalized rank regret is defined in Section 6.1. Human ties, evaluator ties, and incomplete rankings are governed by a frozen policy and reported separately. Pair-only data do not supply missing ranks or cardinal utility. Core H_1 agreement is reported by dimension before any macro average.

Randomness and repeated use. Freeze $R \geq 3$, score precision, retry policy, artifact reuse, and cache identity. Retries are outcome-blind and uniformly applied. Repeated scores are not independent task observations. Paired uncertainty uses original-input/source group resampling, preserving candidate groups and comparisons; the planned interval level is 95%. Report independent search seeds separately from scoring repeats, and distinguish within-run execution variation from between-search variation. Retain failed, rejected, incomplete, and interrupted attempts in the authorized analysis, without favorable resampling or denominator changes.

C2/C3 human outcomes. Report wins, losses, ties, invalid outputs, and major regressions per input group. A prespecified tie-aware preference probability may assign half credit to a tie, with its denominator and missingness bounds stated. Report C2 changes against the shared y_0 in addition to the critic contrast, and C3 changes against Base in addition to the reward contrast. Intervals and costs accompany effect sizes. Multiple domains, dimensions, and ablations require declared primary contrasts and a prespecified multiplicity policy. A search-selected development maximum is not substituted for any final endpoint.

Table 5: Planned ablation conditions and the inference each comparison supports. All conditions require matched evidence access and accounted computation. An oracle condition is diagnostic, not a deployable comparator.

Factor	Conditions	Endpoint or interpretation
Acquisition	Random; judge uncertainty; cross-model conflict; run instability; Semantic multi-source disagreement	Final held-out alignment at equal $N_H \in \{20, 50, 100, 200\}$; acquisition yield alone is insufficient.
Human detail	Same pairs with H_0 , $H_0 + H_1$, $H_0 + H_1 + H_2$, and $H_0 + H_1 + H_2 + H_3$	Marginal value of granularity; report annotation units and annotation time.
Editable scope	Full (A, K, V, T, Π, G) ; rubric/procedure only (K, V) ; fixed tools and static routes; routing only	Whether additional observations or structural changes are needed beyond criteria changes.
Evidence ceiling	Oracle evidence and/or oracle routing with access declared	Diagnostic headroom under privileged information; no deployment claim.
Search allocation	Independent one-shot proposals; history-only proposals; direction depth-first; direction search with pause/resume	Value of direction management at equal total search and measurement budgets.
Tool activation	Tool on/off with compatible consumers	Contribution of acquiring the observation, not the whole program bundle.
Tool implementation	Evolved tool versus seed tool	Contribution of changed measurement or evidence extraction.
Evidence route	Adaptive versus fixed routing with the same tool set	Contribution of which observations reach which modules.
Rubric and modules	Evolved versus fixed criteria/procedures	Contribution of judgment changes with observations held fixed.
Fusion	Declared numerical rule versus semantic synthesis	Contribution of resolving priorities, overlap, and conflicting findings.
Feedback adapter	Direct pairs versus derived ranking/pointwise/error labels on the same subset	Label fidelity and its downstream effect; keep original input dependence.

G FULL ABLATION MATRIX

The comparisons in Table 5 freeze the initial program, evidence access, edit interfaces, proposal opportunities, and measurement budgets. All diagnostic and validation calls are charged, including unsuccessful attempts. If method interfaces differ, report that difference instead of attributing a resource advantage to the search policy. Attribution uses paired original-input groups and pre-specified compatible component replacements. Observed interactions may limit a one-component interpretation.

H PLANNED C3 TRAINING PROTOCOL

C3 is specified to test reward usefulness independently of evaluator agreement and critic usefulness. It is not an executed experiment, and no parameter training is part of this manuscript preparation.

Under the controls in Table 6, the frozen reward produces scalar scores within each rollout group using the same normalization and advantage rule. Zero-variance groups, invalid artifacts, and execution failures require a common prespecified policy. Deterministic checks may enforce format or executability, but do not define the entire reward for open-ended quality.

The pilot uses development validation only. If its prespecified stability or reward-hacking gate fails, expansion stops for that domain and the failure is retained. Success elsewhere is not a substitute. Final sample sizes are determined from pilot variance and a declared minimum detectable effect before test generation. Test inputs and judgments do not participate in training, checkpoint selection, reward design, or failure-driven edits.

Table 6: Controls for the planned LoRA + GRPO comparison. Numerical hyperparameters, pilot thresholds, and final sample sizes remain to be frozen in the experiment manifest.

Control	Contract
Arms	Base is untrained; Seed-GRPO and ERA-GRPO start from the identical base checkpoint.
Training data	Same unlabeled training inputs and order; source-level deduplication and contamination audit.
Policy updates	Same LoRA structure, optimizer, learning rate, update count, group size, KL rule, checkpoint cadence, and policy-training compute.
Online samples	Each updated policy samples its own rollout groups. Do not share trajectories once policies diverge.
Reward version	Freeze program, prompt, tools, routing, aggregation, normalization, ties, abstention, and invalid-output policy before training.
Reward resources	Equal reward-call counts and token caps; report actual inference time, tool use, and monetary or compute cost separately.
Training validation	Input-disjoint development validation selects checkpoints and checks reward exploitation; final metrics and final human labels remain locked.
Pilot gate	Prespecified acceptable validity, reward variance, KL, length, stability, held-out quality, and reward-hacking diagnostics before expansion.
Final generation	Freeze checkpoints and decoding, then generate new Base, Seed-GRPO, and ERA-GRPO outputs on the same unseen ID/OOD inputs.
Final judgment	New blinded H_0 and core H_1 , reserved external metrics, invalid/major-regression rates, length and format shifts, and total cost.

Any deliberate sharing of locked input pools across C1/C2/C3 must be prespecified with an access audit; their newly generated outputs and judgments are never treated as reusable old labels.

I REQUIRED RESULT AND ANALYSIS OUTPUTS

Empirical plots for the outcomes in Table 7 require frozen reporting cohorts and analysis. Suitable outputs include human-budget curves, per-domain component effects, annotated search trajectories, and C2/C3 regression breakdowns. The numerical displays in Appendix J are reporting fixtures, not provisional substitutes for those measurements. Table 8 records the empirical status of the downstream studies independently of the simulated C2 illustration.

Table 7: Required empirical outputs beyond the simulated reporting examples. Every result requires authorized frozen data and a traceable analysis; neither conceptual illustrations nor simulated displays supply such evidence.

Output	Required summary	Interpretation limit
C1 alignment	Global/top-region agreement, per-dimension H_1 , worst stratum, ID/OOD, coverage, repeated-score variation	Pairwise improvement does not establish selection or generation gains.
Feedback quality	Label origin, source protocol, annotator pool, agreement, abstention, unresolved judgments, and annotation time per setting	Newly collected labels are not automatically more reliable.
C1 selection	Best-of- N , normalized rank regret, complete human rankings, grouped intervals	Rank regret is not cardinal utility; pair-only corpora are insufficient.
C2 refinement	New blinded contrast, changes against y_0 , major regressions, dimensional quality, external metrics	Fresh shared initial outputs and matched refinement are necessary.
C3 training	New blinded reward contrast, Base sanity check, external metrics, validity, reward-hacking diagnostics	Design-only until separately authorized training and final evaluation occur.
Efficiency	Human units and time, search calls/tokens/tool use, runtime, deployment and training costs	More compute or annotation detail is not a free method gain.
Trajectory	Directions, implementation attempts, continuation/pause, exhaustion/refutation, resource allocation	A schematic tree is not an optimization trace or evidence of efficiency.
Mechanism	Paired component effects, interactions, remaining gap to human agreement, scoped counterexamples	Assign limitation regimes only after evidence; do not generalize from one stratum.
Adapter fidelity	Direct/derived label agreement and downstream differences on the same subset	Annotation conversion cannot change the construct being evaluated.

Table 8: Empirical downstream status, separate from all simulated displays. Dashes denote unmeasured quantities, not zeros. C3 is a planned study only.

Study	Primary contrast	ID preference	OOD preference	Regressions
C2	ERA-refined vs. Seed-refined	—	—	—
C3	ERA-GRPO vs. Seed-GRPO	—	—	—

J SIMULATED REPORTING FIXTURES

Purpose and limits. All quantitative figures and Table 10 are generated from explicitly simulated data for reviewing the reporting design. The generator does not execute ERA, run a model, consult benchmark outcomes, or recruit annotators. Its assigned noise, effects, costs, and cohort labels are not estimates of expected performance, a power analysis, or evidence for method ranking. Positive, near-zero, and negative contrasts are retained. None of these fixtures provides evidence for acquisition efficiency, component effectiveness, or human-label reliability. No C3 training outcomes are simulated.

Construction and uncertainty. The C1 fixtures cover all 22 domains in Table 1, with 11 illustrative cohorts per feedback setting. Table 9 retains eight fixed tasks for C2 and diagnostic detail; it does not define the study’s scope. A domain may use both feedback settings in an actual study. Gaussian latent utilities define synthetic C1 rankings. Noisy scores are repeated three times per artifact and reduced to medians before strict comparison. Incomplete groups remain in coverage denominators. The C1 table averages domain-level metrics equally within each setting. Its means use each method’s complete groups, whereas the paired differences in Figure 4 use the intersection of completed groups. Intervals use 1,000 input-group bootstrap replicates of these artificial observations. They quantify resampling variation in the fixtures, not uncertainty about ERA. The acquisition

Table 9: Eight diagnostic cohorts for the simulated C2, annotation, and component displays. The C1 displays cover all 22 domains separately. Assignments are layout examples, not verified data availability or fixed domain categories. Each C1 domain/split contains 160 synthetic input groups and four candidates.

Family	Reused-preference example	New-preference example
Structured visuals	TD-SV4: specification-conditioned web	TD-SV2: webpage reconstruction
Images / 3D	TD-IG1: detailed image captioning	TD-IG3: instruction-based image editing
Text generation	TD-TG1: document summarization	TD-TG2: RAG-supported answers
Automated research	TD-AR1: research ideas	TD-AR2: ablation plans

generator assigns budget responses for five policies and 12 toy runs; its bands are percentile ranges across runs, not bootstrap intervals.

Balanced subset in the main figure. Figure 2 uses a subset of the OOD fixtures, with 720 input groups per family. The six-task families contain 120 inputs per task; the five-task families contain 144. Within each task, input identifiers are ranked by SHA256 with fixed seed 20260918, without reference to outcomes. The same selected inputs, including incomplete responses, are retained for every method. Agreement and its 1,000-replicate bootstrap intervals are recomputed on each method’s complete inputs; cost retains all selected inputs. Setting means weight their 11 domains equally, rather than pooling inputs. The ring’s equal quarters reflect this synthetic allocation, not measured domain prevalence or a proposed empirical sampling rule. The released selection list and numerical exports distinguish this display from the full-cohort tables and appendix plots.

Table 10: **Simulated C1 reporting example; not experimental results.** Each setting averages eleven equally weighted toy cohorts. Pair = pairwise agreement; Top = top-region agreement; Bo4 = Best-of-4 accuracy (all in %); Reg = normalized rank regret (lower is better). Cohort assignments do not assert actual label availability. Coverage and domain summaries accompany the reproducible fixtures.

Evaluator	ID test				OOD test			
	Pair	Top	Bo4	Reg	Pair	Top	Bo4	Reg
SIMULATED Reused preferences								
Single metric	66.5	67.5	44.8	0.31	64.1	65.2	43.5	0.32
Metric ensemble	69.9	70.9	51.2	0.26	66.8	68.0	45.3	0.30
Static judge	71.3	72.3	52.2	0.25	69.9	71.1	49.8	0.27
Static tools	74.0	75.2	57.1	0.23	71.3	72.9	52.8	0.24
Prompt optimization	74.1	75.5	56.8	0.22	71.9	72.9	54.4	0.24
Program search	74.8	76.1	57.7	0.21	73.0	74.5	55.5	0.23
ERA	75.3	76.9	61.2	0.20	72.7	74.2	55.3	0.23
SIMULATED New preferences								
Single metric	66.2	67.0	45.1	0.30	63.3	64.3	40.5	0.34
Metric ensemble	68.3	69.6	47.8	0.29	65.9	66.9	43.8	0.32
Static judge	71.3	72.5	53.5	0.24	68.4	69.6	48.0	0.29
Static tools	72.4	73.6	54.9	0.24	70.0	71.2	51.3	0.26
Prompt optimization	74.3	75.5	56.8	0.22	70.6	71.8	51.3	0.26
Program search	74.3	75.7	56.5	0.22	71.6	72.9	53.4	0.24
ERA	75.8	77.1	59.9	0.20	72.2	73.6	53.4	0.24

Expanded alignment and acquisition displays. Figure 4a–b reports per-domain ERA-minus-seed effects with paired intervals for all 22 domains on the full synthetic cohorts. These paired contrasts complement the marginal means for the balanced OOD subset in Figure 2c. Figure 5a–b compares acquisition policies against random sampling at equal human budgets; all policies include the same reliability anchor. Annotation granularity is held fixed here and varied separately in Figure 8b.

SIMULATED DATA

NOT EXPERIMENTAL RESULTS

22 DOMAINS / 4 FAMILIES

Structured visuals · 6

SV1	Brief → single slide	R
SV2	Screenshot → webpage	N
SV3	Data → visualization	R
SV4	Specification → webpage	R
SV5	Sources → slide deck	N
SV6	Topic → slide deck	N

Images / 3D · 6

IG1	Detailed captioning	R
IG2	Text → image	R
IG3	Instruction-based editing	N
IG4	Super-resolution	N
IG5	Text → 3D	R
IG6	Image → 3D	N

Text generation · 5

TG1	Document summarization	R
TG2	RAG long-form answers	N
TG3	Translation	R
TG4	Instruction constraints	N
TG5	Open dialogue	R

Automated research · 5

AR1	Research ideas	R
AR2	Ablation plans	N
AR3	Table → insight	R
AR4	Related work	N
AR5	Peer-review critique	N

a ID test

b OOD test




Figure 4: **Simulated domain-level alignment; not experimental results.** Panels (a) and (b) compare ERA with the static tool-augmented seed on ID and OOD inputs. Diamonds are paired mean differences; whiskers are 95% bootstrap intervals over jointly completed synthetic input groups. Zero denotes no difference. All 22 domains retain the 6/6/5/5 taxonomy order. Colors identify domain families; R/N denote illustrative reused/new feedback assignments, not permanent domain categories or measured data readiness.





Figure 5: **Simulated acquisition curves; not experimental results.** Panels (a) and (b) compare five policies under reused and newly collected preferences. Lines average 12 toy runs; bands mark each policy’s 10th–90th percentiles, not confidence intervals. Budgets count original-input annotation units, including a shared 10-unit random anchor.


Annotation outcomes. The annotation fixture assigns three simulated raters per group. Candidate preferences are encoded as 0/1, a tie as 2, and abstention as -1 . Agreement is computed within each group from all available non-abstaining rater pairs, then averaged across eligible groups. Fewer than two valid votes is insufficient; three distinct valid votes yields no consensus. A two-of-three majority gives a decisive preference or a tie. These outcomes are not evidence about actual annotators or any public benchmark.



Three simulated raters per group. Agreement, abstention, and consensus use distinct denominators.

Figure 6: Simulated annotation audit; no human study was conducted. Panel (a) averages valid-pair agreement within input groups; whiskers are 95% group-bootstrap intervals. Panel (b) divides abstentions by all requested ratings. Panel (c) retains decisive preferences, ties, no consensus, and insufficient votes under a two-of-three rule. Domain codes map to Table 9. The settings share a schema, not an assumed reliability ordering.

Selection beyond pairwise agreement. Figure 7a–b reports the entire selected-rank distribution. Rank 1 is the synthetic human optimum; unscored inputs remain visible. This full-input denominator differs from the completed-group Best-of-4 accuracy in Table 10. Both definitions are exported with their denominators so missingness cannot masquerade as better selection.



ID test; equal-sized toy cohorts. Rank 1 is best. Failed scoring remains in the denominator.

Figure 7: Simulated C1 selection; not experimental results. Panels (a) and (b) show horizontal rank profiles for four evaluator fixtures in the two construction settings, using ID inputs. Rank 1 is best. Labels give its share of all input groups, including unscored groups in the denominator. Colors encode ranks, not methods. Equal cohort sizes permit pooling within a setting; actual domain-level results must also be reported.

Mechanisms and feedback detail. The heatmap in Figure 8a specifies conditional contrasts between a full program and a program holding one component fixed. Such contrasts are not additive causal contributions. Panel (b) varies cumulative feedback on the same synthetic input groups. Its distributions describe assigned group-level values rather than independent search runs. Actual granularity studies must also account for annotation time and unequal label availability.





Figure 8: **Simulated ablations; not experimental results.** Panel (a) gives mean full-minus-fixed agreement (percentage points) for (A, K, V, T, Π, G) across eight diagnostic domains, with cream at zero. Panel (b) colors cumulative H_0 – H_3 levels. Left/right halves denote reused/new preferences; right halves are hatched. Dots are synthetic groups, white markers medians, thick segments quartiles, and thin segments the 10th–90th percentiles. Domain codes follow Table 9. Conditional contrasts are not additive contributions or evidence of component benefits.

Refinement and resource use. The C2 fixture constructs two noisy refinements from a shared latent y_0 . Figure 9a retains their paired quality differences and panel (b) converts those differences to wins, ties, and losses with an independent abstention indicator. The toy tie threshold is 0.15 latent units, not a proposed universal human threshold. Reused/new row labels describe feedback used for evaluator evolution; endpoint judgments on fresh C2 outputs must be new in both settings. Changes against y_0 and major-regression flags are retained in the data for the full reporting contract.

Figure 10a–b aligns assigned relative deployment costs and agreement on common evaluator rows, keeping both settings visible. Connecting segments pair settings, not steps of evolution. A real resource comparison must report search and annotation costs separately from deployment cost; the artificial costs here cannot establish practical efficiency.

Reproduction and replacement. The offline script `scripts/render_simulated_results.py`, seeded with 20260916, produces the CSV fixtures, vector figures, numerical table, and a manifest of hashes. Every data row carries a `SIMULATED` provenance field, and every chart discloses its status within the image. The committed assets compile without model access. Before a results-bearing submission, replace the fixtures with authorized measurements, regenerate captions and uncertainty, and audit all claims against the actual evidence. Removing a disclosure without replacing its data is not a valid conversion to results.



Difference = ERA-refined minus Seed-refined. All artifacts and judgments here are synthetic.

Figure 9: **Simulated C2 outcomes; not experimental results.** Panel (a) combines half-violin densities, all synthetic group values, median markers, and interquartile segments for ERA-refined minus Seed-refined quality. Density colors identify domain families. The horizontal scale is an arbitrary latent unit, not a human score. Panel (b) centers ties between seed and ERA wins, with abstentions in a separate percentage column. Every share uses all input groups under the toy rule. Both construction settings would require fresh blinded human endpoint judgments in an actual C2 study.



Lines pair the two settings for one evaluator. Costs are assigned, not measured compute.

Figure 10: **Simulated quality-cost trade-offs; not measured costs.** Panel (a) gives assigned costs and panel (b) gives agreement, each averaging eleven ID cohorts equally per setting. Colors identify evaluators; metric-only references are gray. Filled circles denote reused preferences and hollow squares new preferences; connecting segments pair the settings for one evaluator. Cost units are assigned to the toy generator, not timed API calls or measured compute. The plot does not establish an efficiency frontier for ERA.

K EXPERIMENT MANIFEST

Table 11: Information required before confirmatory evaluation. Unset values are intentional placeholders, not inferred development defaults.

Item	Required specification
Domain and sources	Versioned taxonomy, decision rule, readiness audit, dataset versions, permissions, input counts, candidate groups: <i>TBD</i>
Supervision	Direct/derived provenance, H_0 – H_3 availability, core dimensions, annotator pool, agreement, time, tie policy: <i>TBD</i>
Partitioning	Original-input/source grouping, deduplication, split sizes, OOD definition, External role, contamination review: <i>TBD</i>
Mining	Source inventory, mappings, dependence groups, anchor size, calibration, strata, frozen policy, human-budget accounting: <i>TBD</i>
Models	Frozen scorer, proposer, tool-model, and generator identities, revisions, decoding, and permissible evidence: <i>TBD</i>
Program and search	Seed, editable components, edit bounds, allocation, development guard, protection criteria, shortlist rule: <i>TBD</i>
Measurement	Repeats $R \geq 3$, precision, completeness, retries, artifact reuse, cache identity, search seeds, stopping rule: <i>TBD</i>
Budgets	Human time, opportunities, model calls/tokens, tools, runtime, deployment cost, training compute, matching rule: <i>TBD</i>
Downstream controls	C1 pool size; C2 fresh generation and refinement; C3 reward, LoRA/GRPO settings, pilot gates and preflight: <i>TBD</i>
Analysis	Primary contrasts, grouped intervals, multiplicity, missingness, reserved metrics, minimum effects, final sample sizes: <i>TBD</i>

Before submission, complete Table 11, replace result placeholders with permitted source-traceable measurements, finalize the empirical scope without changing the method’s definition of a domain, and supply appropriately anonymized reproducibility artifacts. Unmeasured hypotheses remain hypotheses.

L PROGRAM EVOLUTION AND PARAMETER ADAPTATION

Parameter adaptation is a competing approach to learning human evaluation standards, not an assumed failure case. Table 12 distinguishes scorer calibration, evaluator weight fine-tuning, and program evolution by their editable objects. Fine-tuning can learn criteria and tool-use behavior when its training interface permits them. The fixed-structure control below is a declared experimental restriction, not a claim about the limits of fine-tuning in general.

Table 12: Adaptation choices for the planned C1 comparison. Rows specify what changes, not a ranking of effectiveness. All use the same permitted human supervision; observed evidence and total adaptation cost must also be reported.

Approach	Adapted object	What the comparison tests
Score calibration	Combination weights, scale maps, or thresholds over frozen seed signals; model and evidence path fixed	Whether the preference mismatch can be addressed by a different numerical decision rule.
Evaluator fine-tuning	Judge weights from the common base checkpoint; seed adapters, tools, routes, and output interface fixed	Whether learning a new mapping from the same observations to judgments is sufficient.
ERA program evolution	Artifact adapters, criteria, evaluator modules, tools, routing, and fusion; base-model weights fixed	Whether changing the evaluation procedure improves alignment beyond parameter adaptation.

Common supervision and selection. All arms use the same proposal-Train inputs, artifacts, and available human feedback levels. A weight-tuned evaluator learns pointwise scores from the permitted comparisons, with dimensional supervision when available; missing human explanations are not replaced by extra teacher labels in this contrast. Scoring remains pointwise and label-blind. Sealed Train supplies aggregate measurements only, never per-example training targets or gradients. Tuning configurations and checkpoints follow the same terminal-Val and locked-Test/OOD boundary as program candidates. Human units, annotation time, unsuccessful trials, optimization compute, tool calls, and selection costs are all counted. Adaptation and deployment costs are reported separately, with matched-budget comparisons wherever feasible and unmatched resources explicitly identified.

Separating observation from interpretation. The end-to-end comparison permits ERA to change its observation path. It cannot by itself attribute a difference to better interpretation of identical evidence. A shared-observation control therefore freezes (A, T, Π) and supplies the same artifact evidence to fine-tuning and ERA restricted to (K, V, G) . This control, together with the component comparisons in Table 5, tests whether changed observations account for any advantage. Neither lower annotation cost nor superiority to fine-tuning is assumed. A combined weight-and-program arm would be needed to test complementarity.

Inspectable does not mean faithful. An explicit rubric and evidence trace make the stated basis of a judgment available for inspection. They do not establish that the judgment follows that basis or matches the human standard. Mechanism assessment must check cited observations and test predicted decision changes under controlled evidence or component interventions. These tests are distinct from preference agreement.

The tuning comparison concerns the *evaluator*. C3 instead trains a *generator* against frozen evaluator rewards. Neither experiment has been executed in this draft. The simulated C1 displays illustrate a subset of reporting conditions and contain no fine-tuning or calibration outcomes.